


(a) Conventional path


2 PCIe crossings · frame staged in host RAM



Plaintext + keys exposed to the host; secure preprocessing steals GPU SM capacity + memory bandwidth from the serving model.

(b) DPU-terminated path

1 PCIe crossing · only the tensor crosses



Plaintext + keys confined to the DPU; GPU SM capacity + memory bandwidth are preserved entirely for inference.